

CSA1223 — Computer Architecture for Multicore Systems

Assessment Tool 2 — Debugging & Optimisation Task | CO2 | Model Answers

Question 1 — Overflow in Signed 2's-Complement Addition

Problem: Adding +120 and +50 as 8-bit signed (2's-complement) numbers gives a wrong result. We need to find why, and how to detect the error in hardware/software.

Step 1 — Binary Computation

Value	8-bit Binary	Decimal
+120	0111 1000	120
+50	0011 0010	50
Sum	1010 1010	-86 (as read)

Mathematically, $120 + 50 = 170$. But an 8-bit signed register can only hold values from -128 to +127. Since 170 exceeds +127, the true sum cannot be represented — the register wraps around and is misread as a negative number (-86).

Step 2 — Identifying the Bug (Root Cause)

- Two positive operands produced a negative result — this alone signals overflow.
- Formal check — track the carry into and out of the sign bit (bit 7): Carry-in to bit 7 = 1, Carry-out of bit 7 = 0
- Since Carry-in $\neq$ Carry-out at the sign bit, the addition has overflowed the representable range.

Step 3 — Proposed Solution

- Hardware overflow rule: Overflow (V) = Carry-in(MSB) XOR Carry-out(MSB)
- Simple sign rule: if both operands share the same sign and the result's sign differs from theirs, overflow has occurred (set the V/overflow flag).
- Practical fixes: perform the addition in a wider register (16-bit/32-bit) when the range is uncertain, check the overflow flag before trusting the result, or use saturating arithmetic that clamps to +127/-128 in embedded systems instead of wrapping silently.

Question 2 — Ripple Carry Adder Delay → Carry Look-Ahead Adder

Problem: A ripple carry adder (RCA) causes noticeable delay in a real-time system. We need to trace the bottleneck and optimise it.

Step 1 — Debugging the Bottleneck

- In an n-bit RCA, each full adder's Carry-Out feeds directly into the Carry-In of the next stage.
- The most-significant sum bit cannot be finalised until the carry has rippled sequentially through every earlier stage.

- This creates a worst-case critical path that grows linearly with word size — roughly proportional to n (2 gate delays per stage) — which dominates total delay for wide (32-bit / 64-bit) operands.

Step 2 — Optimised Solution: Carry Look-Ahead Adder (CLA)

A CLA computes every carry directly from the inputs, in parallel, instead of waiting for it to ripple through:

$$\begin{aligned} \text{Generate: } G_i &= A_i \cdot B_i & \text{Propagate: } P_i &= A_i \oplus B_i \\ C_1 &= G_0 + P_0C_0 \\ C_2 &= G_1 + P_1G_0 + P_1P_0C_0 \\ C_3 &= G_2 + P_2G_1 + P_2P_1G_0 + P_2P_1P_0C_0 \end{aligned}$$

- Each carry is a direct Boolean function of the inputs and initial carry, so it can be evaluated in about 2 gate delays instead of waiting n stages.
- Cascading look-ahead blocks brings overall delay down from $O(n)$ to roughly $O(\log n)$, at the cost of more gates and higher fan-in — an acceptable trade-off when speed matters more than gate count.

Comparison

Aspect	Ripple Carry Adder	Carry Look-Ahead Adder
Carry generation	Sequential (stage by stage)	Parallel (direct Boolean expression)
Delay	$O(n)$	$O(\log n)$
Hardware cost	Low (simple, few gates)	Higher (more gates, larger fan-in)
Best suited for	Small, low-cost designs	Real-time / high-speed systems

Question 3 — Booth's Algorithm — Errors with Negative Operands

Problem: A Booth's-algorithm multiplier gives wrong answers for certain negative operands. We trace bit-pair checking, shifting and sign handling to find the bug.

Step 1 — Common Sources of the Bug

- Using a logical shift right (fills with 0) instead of an arithmetic shift right (ASR, which copies/sign-extends the MSB) when shifting the $\{A, Q, Q-1\}$ register — this destroys the sign for negative partial results.
- Forgetting to initialise the extra bit $Q-1$ to 0 before the first cycle.
- Implementing 'subtract' incorrectly, instead of as Add of the 2's-complement of the multiplicand M .
- Running for the wrong number of cycles — it must run for exactly n cycles, where n is the operand width.

Step 2 — The Correct Bit-Pair Rule

$Q_n \ Q_{-1}$	Operation	Then
0 0	No arithmetic operation	Arithmetic Shift Right (ASR)

Qn Q-1	Operation	Then
0 1	$A = A + M$	ASR
1 0	$A = A - M$	ASR
1 1	No arithmetic operation	ASR

Step 3 — Corrected Worked Example ($M = -3$, Multiplier = $+5$, $n = 4$ bits)

$M = 1101$ (-3), $-M = 0011$, $Q = 0101$ ($+5$). Every shift below correctly sign-extends the MSB of A.

Step	Qn Q-1	Operation	A	Q	Q-1
Init	—	—	0000	0101	0
1	1 0	$A = A - M$	0001	1010	1
2	0 1	$A = A + M$	1111	0101	0
3	1 0	$A = A - M$	0001	0010	1
4	0 1	$A = A + M$	1111	0001	0

Final product $\{A, Q\} = 1111\ 0001 \rightarrow$ this 8-bit two's-complement value equals -15 , which matches $-3 \times 5 = -15$ exactly.

Step 4 — Correction Applied

- Always sign-extend the MSB of A during every ASR (never zero-fill).
- Initialise $Q-1 = 0$ before the first cycle.
- Perform subtraction as addition of the 2's-complement of M.
- Run for exactly n cycles matching the operand width.

Question 4 — Restoring Division Inefficiency $\rightarrow$ Non-Restoring Division

Problem: An embedded system's restoring-division routine is slow because of repeated restoration steps. We trace the workflow and optimise it.

Step 1 — Debugging Restoring Division

- Each iteration: shift the remainder left, then subtract the divisor ($R = R - D$).
- If R becomes negative, the hardware must restore it by adding the divisor back ($R = R + D$) and record quotient bit 0; otherwise it keeps R and records quotient bit 1.
- That extra 'restore' addition is wasted work — it happens whenever the trial subtraction goes negative, which is roughly half of all iterations on average — inflating the worst case to about $2n$ arithmetic operations for an n -bit division. This is exactly the inefficiency causing the slowdown.

Step 2 — Optimised Solution: Non-Restoring Division

- Never explicitly restores. Instead, the sign of the previous remainder decides the next operation:
 - If the previous remainder was non-negative $\rightarrow$ subtract the shifted divisor
 - If the previous remainder was negative $\rightarrow$ add the shifted divisor

- Quotient bit = 1 if the resulting remainder is non-negative, else 0.
- Exactly one add/subtract per iteration (n operations total), plus a single final correction step (add the divisor back once, only if the last remainder is negative) → n + 1 operations in total, versus up to 2n for restoring division.

Comparison

Aspect	Restoring Division	Non-Restoring Division
Operation per step	Subtract, then restore (add) if negative	Subtract OR add, based on previous sign
Extra restore step	Yes — up to n times	No — only one final correction if needed
Worst-case operations	~2n	n + 1
Timing	Less predictable (variable restores)	Predictable, fixed cycle count

Question 5 — IEEE 754 Precision Loss: Very Small + Very Large Numbers

Problem: Adding a very large and a very small IEEE 754 single-precision float gives an inaccurate result. We examine exponent alignment, normalisation and rounding.

Step 1 — IEEE 754 Single-Precision Format (32-bit)

Sign	Exponent	Mantissa (Fraction)
1 bit	8 bits (bias = 127)	23 bits (+ implicit leading 1)

Step 2 — Debugging the Loss of Accuracy

- Before adding, the two exponents must be equalised: the mantissa of the smaller-magnitude number is shifted right by ($E_{\text{large}} - E_{\text{small}}$) positions.
- If that exponent difference exceeds about 24 (23 mantissa bits + 1 implicit bit), the smaller number's mantissa is shifted completely out of the register — it contributes nothing to the sum at all.
- Example: adding 1.0×10^{10} and 1.0×10^{-10} in single precision returns exactly 1.0×10^{10} — the small term is silently swallowed.
- This is a genuine loss-of-significance / rounding error inherent to limited mantissa width, not a coding mistake — but it is worsened if no guard/round/sticky bits are kept during the shift, since bits are then discarded before rounding even gets a chance to compensate.

Step 3 — Optimisation Techniques

- Guard, Round, Sticky (GRS) bits: retain a few extra bits during alignment so the final round-to-nearest-even decision uses the true discarded bit pattern instead of a blind truncation.
- Use double precision (64-bit: 11 exponent bits, 52 mantissa bits) for critical or intermediate calculations — the larger mantissa delays (though does not eliminate) the point at which small terms vanish.
- Kahan summation algorithm: keeps a running compensation value that captures the low-order bits lost on each addition, giving far better accuracy over long summations without changing data type.

- Reorder additions from smallest magnitude to largest (pairwise summation) so values of comparable size combine first, minimising the magnitude gap at every step.